UNIT IV

MapReduce Workflow,MRUnit, Testing Data, Local Test, Anatomy of MapReduce Job Run Classic MapReduce vs YARN, Failures in MapReduce & YARN, Job Scheduling Shuffle & Sort, task Execution, Input & Output Formats, Map Reduce Types

MapReduce Workflow:

MapReduce is a Hadoop processing framework that efficiently handles large-scale data across distributed machines. Unlike traditional systems, it works directly on data stored across nodes in HDFS.
Hadoop MapReduce follows a simple yet powerful data processing model that breaks large datasets into smaller chunks and processes them in parallel across a cluster. This flow from input splitting to mapping, shuffling and reducing ensures scalable, fault-tolerant and efficient data processing over distributed systems.
Below is the workflow of Hadoop MapReduce with a simple data flow diagram.

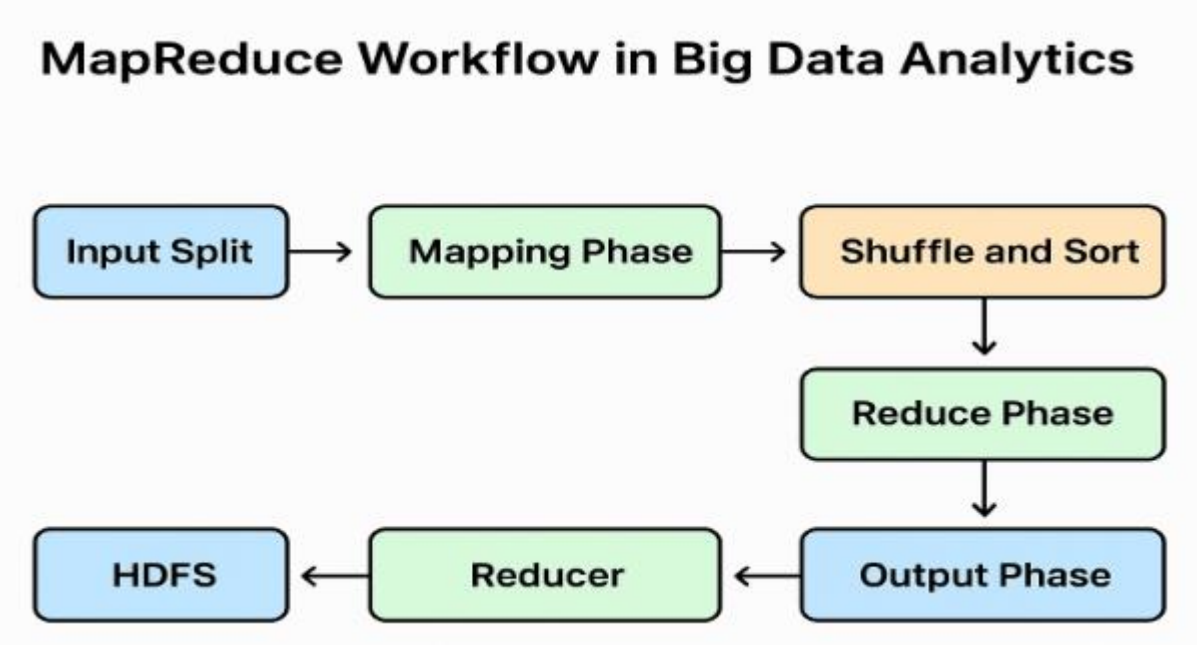

1. Input Split
The input data (e.g., a big log file or dataset) is divided into smaller chunks called Input Splits. Each split is processed independently by a separate Mapper.
Example: If you have a 1 GB file and Hadoop splits it into four 256 MB chunks, it will use 4 Mappers one for each chunk.
2. Mapper Phase
Each Mapper runs in parallel on different nodes and processes one input split.
What it does:
- Reads the input data line by line
- Transforms each line into key-value pairs

- Stores the intermediate output locally (not yet in HDFS)

Example: If the task is counting words, the Mapper reads: "Data is power"--> emits:
("Data", 1), ("is", 1), ("power", 1)

3. Shuffling & Sorting

This is a behind-the-scenes phase handled by Hadoop after mapping is done.
What it does:

- Shuffles intermediate key-value pairs across the cluster
- Groups all values with the same key
- Sorts them by key before sending to the Reducer

Example: From all Mappers, these pairs:
("Data", 1), ("Data", 1), ("power", 1)
are grouped into:
("Data", [1, 1]), ("power", [1])

4. Reducer Phase

Each Reducer receives a list of values for each unique key.
What it does:

- Applies aggregation logic (e.g., sum, average, filter)
- Generates the final key-value output
- Stores the result in HDFS

Example: For word count, the Reducer gets:
("Data", [1, 1]) --> outputs: ("Data", 2)
The final output is saved in files like: part-r-00000

5. Output Phase

Final results are written back to **HDFS**.Each Reducer writes its own output file (part files).

MR Unit

MRUnit is a **Java-based unit testing framework** specifically designed for testing Hadoop **MapReduce** programs. Just like JUnit is used to test general Java programs, MRUnit provides tools to test the correctness of **Mapper**, **Reducer**, and **MapReduce** jobs without having to run a full Hadoop cluster.

In a large distributed environment like Hadoop, debugging is difficult because jobs run across many machines, and errors may be hard to trace. MRUnit solves this by allowing developers to test their logic **locally, quickly, and deterministically**. It helps ensure that MapReduce components behave exactly as expected before deploying them to a cluster.

Why MRUnit is Needed?

MapReduce jobs often process massive datasets. Running the job on a real cluster for every code change is time-consuming and impractical. Developers need a way to verify:

- Whether Mapper emits correct key-value pairs
- Whether Reducer aggregates values correctly

- Whether the overall flow (Map → Shuffle → Reduce) is correct
- Whether corner cases, invalid data, or special patterns are handled

MRUnit allows this by **mocking the Hadoop execution environment**. It simulates the behavior of the framework's map and reduce phases so you can focus only on logic testing.

The main goal is **early error detection**, which saves time, cluster resources, and avoids failures in production.

MRUnit Architecture & Components

MRUnit provides three main driver classes that mimic the actual MapReduce pipeline:

1) MapDriver

Used to test only the **Mapper** logic.

It allows developers to:

- Provide custom input key-value pairs
- Run the mapper
- Verify the exact output emitted by the mapper

2) ReduceDriver

Used to test only the **Reducer** logic.

It helps validate:

- Whether reducer receives grouped values properly
- Whether the output aggregation is correct

3) MapReduceDriver

Used to test the **entire pipeline** (Mapper → Shuffle → Reducer).

It automatically simulates:

- Mapping
- Grouping of intermediate keys
- Reducing

This ensures the final result is correct from end to end.

How MRUnit Works Internally

MRUnit does **not** need Hadoop to run. Instead, it creates lightweight objects that mimic Hadoop's:

- Context objects
- Writable data types
- Key-value structures
- MapReduce pipeline

During a test, MRUnit performs the following:

1. Takes input fed by the tester.
2. Invokes the map function and captures outputs.
3. For reduce tests, it groups inputs as Hadoop would.
4. Invokes reducer with grouped values.
5. Compares produced output against expected output.
6. Reports success or failure, just like JUnit.

MRUnit uses assertions internally, making testing highly reliable.

Advantages of MRUnit

1. Faster Development

Helps developers catch bugs early without using expensive cluster time.

2. Local Testing

Allows execution locally on a developer machine without HDFS or Hadoop daemons.

3. Accurate Validation

Tests exact mapper/reducer output, ensuring logical correctness.

4. Easy to Integrate with JUnit

Allows both frameworks to work together, improving testing quality

Testing Data

Testing data in Big Data Analytics (BDA) refers to the process of validating the accuracy, quality, performance, and reliability of big data systems and analytical outcomes. Since big data is characterized by high **volume, velocity, variety, veracity, and value**, testing becomes essential to ensure that the data processing pipeline works correctly at scale.

1. Need for Testing in BDA:

Big data systems involve distributed storage (HDFS, NoSQL), large-scale processing (Spark, MapReduce), and real-time ingestion (Kafka, Flume). Testing ensures data accuracy, completeness, consistency, scalability, fault tolerance, and correctness of analytics and machine learning outputs.

2. Types of Testing in BDA:

- **(i) Data Quality Testing:** Verifies accuracy, completeness, validity, timeliness, uniqueness, and consistency of data.
- **(ii) Data Ingestion Testing:** Ensures correct extraction and loading using tools like Kafka, Sqoop, NiFi; checks schema validation, data loss, ordering, and ingestion speed.
- **(iii) ETL / Transformation Testing:** Validates business rules, aggregations, joins, filtering, and data mapping in Spark/MapReduce. Reconciliation between source and output is performed.
- **(iv) Pipeline (End-to-End) Testing:** Tests the complete workflow from ingestion → storage → processing → analytics to verify logic, completeness, and system reliability.
- **(v) Performance & Load Testing:** Measures throughput, cluster scalability, latency, resource usage, and behavior under large data loads using JMeter, Gatling, etc.
- **(vi) Security Testing:** Checks authentication, authorization, encryption, and auditing mechanisms in big data environments.

3. Big Data Testing Process:

Testing typically follows a layered approach:

1. **Ingestion Layer Testing** (data extraction and loading),
2. **Storage Layer Testing** (HDFS/Hive consistency, partitioning, replication),
3. **Processing Layer Testing** (Spark/MapReduce job accuracy and performance),
4. **Analytics Layer Testing** (BI reports, ML model validation).

4. Test Data Generation Techniques:
Test data may be generated using synthetic tools, sampling real datasets, injecting noise/outliers, creating extreme cases, and simulating high-velocity streams to check system robustness.

5. Challenges in Big Data Testing:
Challenges include handling extremely large datasets, diverse data formats, distributed architecture, high ingestion speed, maintaining test environments, and difficulty in debugging failures across multiple nodes.

6. Best Practices:
Use automation, validate at every stage, monitor logs, implement schema validation, test at production scale, use reusable test datasets, ensure security compliance, and conduct regular regression tests for pipelines.

Conclusion:
Testing in Big Data Analytics ensures that the entire data ecosystem—ingestion, storage, processing, and analytics—operates correctly, efficiently, and securely. It is crucial for generating trustworthy insights and supporting large-scale data-driven decision-making.

Anatomy of a MapReduce Job Run

A MapReduce job run involves multiple coordinated stages that begin from job submission and end with producing output in HDFS. These steps ensure scalability, fault tolerance, parallelism, and reliability in Hadoop's distributed computing environment.

The complete lifecycle consists of the following major phases:

1. **Job Submission**
2. **Job Initialization**
3. **Task Assignment**
4. **Map Task Execution**
5. **Map-Side Shuffle and Spill**
6. **Reduce-Side Shuffle and Sort**
7. **Reduce Task Execution**
8. **Job Completion**

1. Job Submission

When a user submits a job using `hadoop jar ...`, the client interacts with the Hadoop cluster to prepare and launch the job.

1.1 Steps during Job Submission

- The driver program (client) sets:
  - Mapper class
  - Reducer class
  - Input/Output formats
  - Input/Output paths
- The client connects to:
  - **JobTracker** in Classic MapReduce (MRv1), OR
  - **ResourceManager** in YARN (MRv2)

1.2 Job Resources

Before running, Hadoop copies job-related data into a **job staging directory** on HDFS:

/tmp/hadoop-yarn/staging/<user>/.staging/job_XXXX/
This directory contains:

- Job JAR file
- Job configuration XML
- Input split metadata
- Shell scripts and libraries

2. Job Initialization

In this phase, the master node (JobTracker/RM) prepares job structures and determines the tasks that need to be executed.

2.1 Input Splits Creation

- The InputFormat class (commonly TextInputFormat) divides the input data into logical chunks called **InputSplits**.
- Each split usually equals the HDFS block size (128 MB default).
- Each InputSplit results in **one map task**.

2.2 JobTracker/ResourceManager Initializes Task Metadata

- Creates a job object.
- Determines:
  - Number of map tasks = number of splits
  - Number of reduce tasks (user-defined)
- Prepares job scheduling information.

3. Task Assignment

Once the job is initialized, the cluster begins allocating nodes to run the map and reduce tasks.

3.1 Task Scheduling

- Hadoop attempts to place **map tasks near the data** (data locality).
- Three locality levels:
  - **Node-local** → Best
  - **Rack-local**
  - **Off-rack** → Worst

3.2 Communication

- JobTracker/ResourceManager assigns tasks.
- TaskTracker/NodeManager launches tasks in JVMs.

4. Map Task Execution

A map task transforms raw input data into intermediate key-value pairs.

4.1 Steps of Map Task

1. **InputSplit assigned to mapper**
2. **RecordReader** reads raw data and generates key/value pairs for the mapper.
3. **map() function** processes each record.
4. Output is collected in an in-memory buffer.

4.2 Mapper Data Flow Diagram

```
InputSplit
    |
    v
RecordReader
    |
    v
map(key, value)  → intermediate (k2, v2)
    |
    v
Map Output Buffer (in memory)
```

5. Map-Side Shuffle & Spill

Shuffle is the process of **transferring map output to reducers**, but before this transfer happens, each mapper performs sorting and partitioning.

5.1 In-Memory Buffering

- Map outputs are stored in a circular memory buffer (default 100 MB).
- When the buffer reaches 80% full → SPILL event is triggered.

5.2 Spill Process

During spill:

- Data is **partitioned** (one partition per reducer).
- Within each partition, data is **sorted by key**.

- **Combiner** may run (optional) to reduce data size.
- Sorted partitions are written to **spill files** on disk.

5.3 Final Merge

- All spill files are merged into a single **Map Output File per partition**.

6. Reduce-Side Shuffle and Sort

This phase is the **core of MapReduce** and often the most time-consuming.

6.1 Shuffle Phase

- Reducers **fetch map output** from all map nodes over HTTP.
- For R reducers and M mappers:
  - Each reducer fetches **M partitions**
  - Total partitions = M × R

6.2 Sort Phase

- Reducers merge all fetched segments.
- The merged data is **sorted by key**.
- Groups values belonging to the same key.

7. Reduce Task Execution

Once data is shuffled and sorted, the reduce phase begins.

7.1 Steps

1. For each unique key, reducer gets:
2. key → list of all values
3. reduce() aggregates or processes the values.
4. Reducer writes the final output to HDFS using OutputFormat.

8. Job Completion

- As reducers finish, they report status to JobTracker/ResourceManager.
- When all tasks succeed:
  - Master marks the job as **SUCCESS**.
  - Temporary staging directories are cleaned.

Full Lifecycle Diagram

```
MAPREDUCE JOB RUN
                  ------------------

1. Job Submission
     |
     v
2. Job Initialization
     |
     v
3. Task Assignment
     |
     v
4. MAP PHASE
   - Read Split
   - Map()
   - Spill
     |
     v
5. Shuffle (Map Side)
     |
     v
6. SHUFFLE (Reduce Side)
   - Fetch
   - Merge
   - Sort
     |
     v
7. REDUCE PHASE
   - Reduce()
   - Write Output
     |
     v
8. Job Completion
```

Classic MapReduce vs YARN

Category	Classic MapReduce (MRv1)	YARN (MRv2)
Core Architecture	Monolithic architecture using **JobTracker + TaskTrackers**.	Modular and distributed using **ResourceManager + NodeManagers + ApplicationMasters**.
Resource Manager Role	**JobTracker** handles both resource management and job scheduling.	**ResourceManager (RM)** only manages cluster resources; job execution is delegated to ApplicationMasters.
Application Execution Manager	None; JobTracker oversees all job lifecycles directly.	**ApplicationMaster (AM)** created per application; manages the job from submission to completion.
Node-level Agent	**TaskTracker** runs map and reduce tasks.	**NodeManager (NM)** provides generalized container execution (not tied to MR tasks).
Task Execution Model	Fixed slots: **map slots** and **reduce slots** permanently predefined.	Dynamic **containers** allocated based on memory and CPU, no fixed map/reduce slots.
Scalability	Limited to ~4,000 nodes due to JobTracker overhead.	Scales to >10,000 nodes because duties are distributed.
Bottlenecks	Single JobTracker becom	

Failures in MapReduce & YARN

1. Failures in Classic MapReduce (MRv1)

Key Components

- **JobTracker (JT)** → central controller
- **TaskTracker (TT)** → worker nodes running map/reduce tasks

Types of Failures

Failure Type	Behavior in MRv1	Recovery Mechanism
Task Failure	Map or Reduce task fails due to JVM crash, hardware issue, or application error.	JT detects failure via heartbeat. Task is retried on same or different TT (default 4 retries).
TaskTracker Failure	TT stops sending heartbeats to JT (node goes down).	All running tasks on that TT are rescheduled to other TTs. JT keeps track of failed tasks.
JobTracker Failure	JT fails → entire cluster loses job tracking and resource management.	Recovery is difficult; older MRv1 versions could not recover jobs (manual restart needed). Later, **JobTracker HA** was introduced but rarely used.
Speculative Execution	Optional: duplicate slow-running tasks to reduce stragglers.	Helps mitigate performance issues but does not prevent failures.
Data Loss	If HDFS block is unavailable, tasks fail.	HDFS replication handles block recovery; MR retries task on node with valid block.

Characteristics

- MRv1 is **single-point-of-failure (JobTracker)**.
- Task failures are handled **at task granularity**, but JobTracker failure is catastrophic.
- Resource utilization is low; idle slots may exist while tasks fail and retry.

2. Failures in YARN (MRv2)

Key Components

- **ResourceManager (RM)** → global resource manager
- **NodeManager (NM)** → node agent
- **ApplicationMaster (AM)** → per-application controller

Types of Failures

Failure Type	Behavior in YARN	Recovery Mechanism
Container/Task Failure	Task running inside a container fails.	AM resubmits task in a new container on same or different node.
NodeManager Failure	Node goes down or NM fails.	RM marks node unhealthy; AM reschedules all containers previously on that node.
ApplicationMaster Failure	AM crashes during job execution.	RM restarts AM for the application. Job continues from last checkpoint (some frameworks handle state).
ResourceManager Failure	RM fails (rare in HA setups).	HA-enabled RM allows **active-standby** RM failover with no job loss.
Speculative Execution / Stragglers	AM can launch duplicate containers for slow tasks.	Helps maintain job performance.
Data Loss	Handled by HDFS replication.	Tasks retried on nodes with valid data.

Characteristics

- **No single point of failure** (HA RM, per-app AM).
- Containerized execution allows **fine-grained task recovery**.
- Better **resource utilization** during failures; tasks rescheduled dynamically.
- Supports **multi-framework failure recovery** (Spark, Tez, Flink) through AM logic.

3. Key Differences: Failure Handling

Feature	Classic MapReduce (MRv1)	YARN (MRv2)
Single Point of Failure	JobTracker → cluster-wide impact	No single point; RM HA + AM per app
Task/Container Recovery	Retry failed task on other TT	AM reschedules failed container dynamically
Job Recovery	Difficult if JobTracker fails	AM restart allows job to continue; RM HA handles RM failures
Node Failure Handling	TT failure → JT reschedules tasks	NM failure → AM reschedules containers on healthy nodes
Speculative Execution	Optional, task-level	Optional, container-level, framework-aware
Overall Reliability	Moderate, limited by JT	High, robust for large clusters and multi-framework workloads

Job Scheduling
Shuffle & Sort

Job Scheduling

Job scheduling in Big Data ensures that multiple jobs are executed efficiently on a cluster, making optimal use of resources and minimizing delays. In Hadoop, jobs are split into **map tasks** and **reduce tasks**, which run on different nodes.

In Hadoop 1, the **JobTracker** schedules tasks, whereas in Hadoop 2 (YARN), the **ResourceManager** allocates cluster resources, and the **ApplicationMaster** manages job execution. Scheduling policies determine **the order and allocation of resources**:

- **FIFO (First In, First Out):** Jobs run in order of submission; simple but can block small jobs.
- **Fair Scheduler:** Allocates resources evenly over time; prevents starvation.
- **Capacity Scheduler:** Divides cluster resources into queues with guaranteed capacity; ideal for multi-tenant environments.

Scheduling also considers **data locality**, meaning tasks are ideally run on nodes where the required data resides, reducing network overhead.

Shuffle & Sort

Shuffle and Sort happens **between Map and Reduce phases**. After mappers generate key-value pairs, the system must **group all values by key** and send them to the correct reducer. This process is critical for **correct aggregation**.

- **Shuffle:** Transfers mapper output to reducers based on the key, using a partitioner (e.g., hash(key) % numReducers).
- **Sort:** Ensures keys are ordered for reducers.

Example: Mapper output: (apple,1), (banana,1), (apple,1)
After Shuffle & Sort:

Reducer 0: (apple,[1,1])
Reducer 1: (banana,[1])

Optimization: Combiners can pre-aggregate data on the mapper side, reducing network transfer during shuffle. Shuffle & Sort is network-intensive and directly affects job performance.

Linking Scheduling and Shuffle & Sort

Effective **job scheduling** ensures tasks run where resources and data are available, reducing shuffle overhead. Meanwhile, **shuffle and sort** organizes data so reducers can efficiently produce final results. Both work together to make MapReduce jobs **fast, scalable, and fault-tolerant**.

Task execution in big data analytics is the process of a distributed computing framework breaking down a large job into smaller tasks that are executed in parallel across multiple nodes in a cluster. This process involves job submission, task assignment by a scheduler, parallel execution on data nodes, sorting and shuffling of intermediate data, and final output generation, all managed by the framework to ensure scalability, fault tolerance, and efficiency.

Key stages of task execution

- **Job submission:**
  A user submits a big data job, such as a MapReduce job, to the cluster.
- **Job initialization:**
  The framework initializes the job by setting up configurations and creating necessary files.
- **Task assignment:**
  A resource manager (like YARN's scheduler) assigns tasks to specific containers on different nodes, considering factors like data locality.
- **Data locality:**
  The system tries to execute a task on the node where its data resides to minimize network traffic. If this isn't possible, the task runs in remote mode, fetching data from another server.

- **Parallel execution:**
  Tasks are executed in parallel. For example, in MapReduce:
- **Mapper execution:** "Mapper" tasks read input data, process it, and produce intermediate key-value pairs.
- **Sort and shuffle:** The intermediate data is sorted and partitioned by key before being sent to reducers.
- **Reducer execution:** "Reducer" tasks process the sorted, intermediate data to produce the final output.
- **Output and commit:**
  Results are first written to a temporary location. Once a task completes successfully, its output is moved to the final output path.
- **Fault tolerance:**
  The framework handles node or task failures by rerunning failed tasks. If speculative execution is used, it ensures only one copy of the output is committed.
- **Job completion:**
  The job is considered complete when all tasks are finished, and the framework reports the status.

input and output formats

input and output formats in big data analytics are crucial for performance and include standard text-based formats like CSV and JSON, as well as specialized binary formats like **Avro**, **Parquet**, and **ORC**. Input formats define how data is read into the system, while output formats dictate how results are written, with options ranging from simple text files to structured binary files and database tables. The choice of format depends on data structure, processing needs, and performance requirements.

Input formats

- **TextFileInputFormat:** The standard, treating each line of a file as a separate record. It's suitable for raw, line-based data like log files.

- **TextInputFormat (MapReduce):** A default text input format that reads files line by line. The key is the byte offset of the line, and the value is the line's content.

- **SequenceFileInputFormat:** Used for reading binary sequences of data, which can be useful as input for another MapReduce job.

- **DBOutputFormat:** Handles input and output for relational databases and HBase.

Output formats

- **TextOutputFormat:** A default output format that writes records as plain text, typically separating key and value with a tab character.
- **SequenceFileOutputFormat:** Writes a sequence of binary output to a file, useful for subsequent MapReduce processes.
- **DBOutputFormat:** Writes results to relational databases or HBase tables.

Other common formats

- **CSV:** Comma-separated values, a widely used but not always optimal format for large datasets.
- **JSON:** JavaScript Object Notation, commonly used for online applications but can be slower for big data processing because it's not highly typed or schema-enriched.
- **Avro:** A row-based data serialization system with a rich data structure and schema.
- **Parquet:** A columnar storage format that is highly efficient for big data analytics due to its compression and encoding capabilities.
- **ORC:** Another columnar format optimized for Hive, offering high performance.

Choosing the right format

- **Data structure:** Consider if your data is structured, semi-structured, or unstructured, and if it needs to support nested data.
- **Performance:** Columnar formats like Parquet and ORC are often faster for queries because they only read the columns needed.
- **Use case:** For simple text files, a text format is fine. For complex, large-scale analytics, specialized formats like Parquet or Avro are better suited.
- **Schema evolution:** Formats with a schema, like Avro and Parquet, can be better for managing changes in data structure over time.

MapReduce Types in Big Data

MapReduce jobs can be classified based on **how many map and reduce phases they have**, or based on **their purpose/operation**. Understanding the types helps in designing jobs effectively.

1. Based on Number of Phases

a) Map-Only Jobs

- **Definition:** Jobs that consist of only **map tasks** without any reducers.
- **Use Case:** When no aggregation is needed; the output of each map task is final.
- **Example:** Filtering records, extracting specific fields from logs.

Flow:

Input → Mapper → Output

b) Map-Reduce Jobs (Single Reduce)

- **Definition:** Standard MapReduce job with **map phase → shuffle & sort → reduce phase**.
- **Use Case:** Most common; used when aggregation or summarization is required.
- **Example:** Word Count, Summing Sales per Product.

Flow:

Input → Mapper → Shuffle & Sort → Reducer → Output

c) Multi-Stage MapReduce

- **Definition:** Jobs that consist of **multiple map-reduce stages** chained together, where output of one stage becomes input for the next.
- **Use Case:** Complex workflows requiring multiple transformations or aggregations.
- **Example:** Sorting and counting, iterative algorithms like PageRank.

Flow:

Input → Mapper1 → Shuffle & Sort → Reducer1 → Mapper2 → Shuffle & Sort → Reducer2 → Output

2. Based on Operation/Functionality

Type	Description	Example
Filtering	Extracts specific data from input	Extract logs with error=500
Aggregation / Summarization	Combines values by key	Total sales per region
Join	Combines datasets based on key	Join customer info with orders
Sorting	Orders data based on key	Sort users by ID or date
Inverted Index	Creates a word → document mapping	Search engines indexing